


Model Comparison — Summary Metrics


Model	Accuracy	Balanced Accuracy	Precision (macro)	Recall (macro)	F1 (macro)	MCC	AUC (macro, OVR)	Avg. Loss	Baseline Acc.
ConvNeXt	0.9633	0.9623	0.9634	0.9623	0.9622	0.9593	0.9991	0.1180	0.1111
EfficientNetV2	0.9652	0.9630	0.9679	0.9630	0.9648	0.9614	0.9995	0.0997	0.1111
ResNeXt	0.9615	0.9607	0.9624	0.9607	0.9606	0.9573	0.9989	0.1168	0.1111

Per-Metric Comparison (green = best, red = worst)


Accuracy




Balanced Accuracy




Precision (macro)




Recall (macro)




F1 (macro)



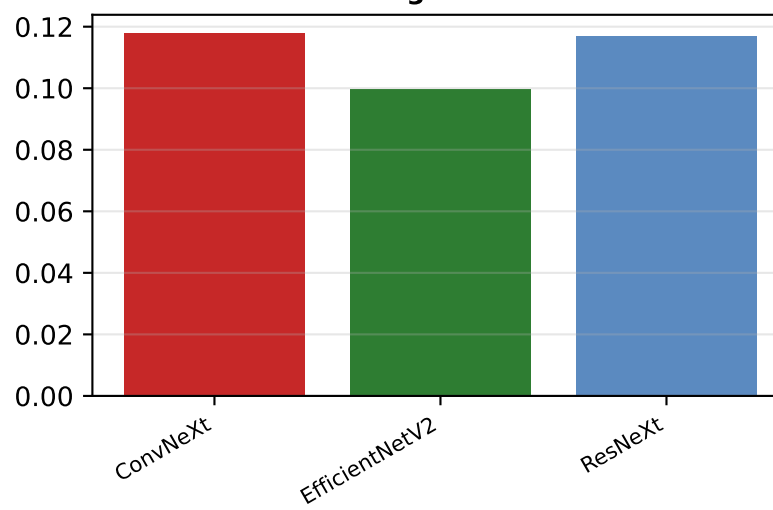
MCC



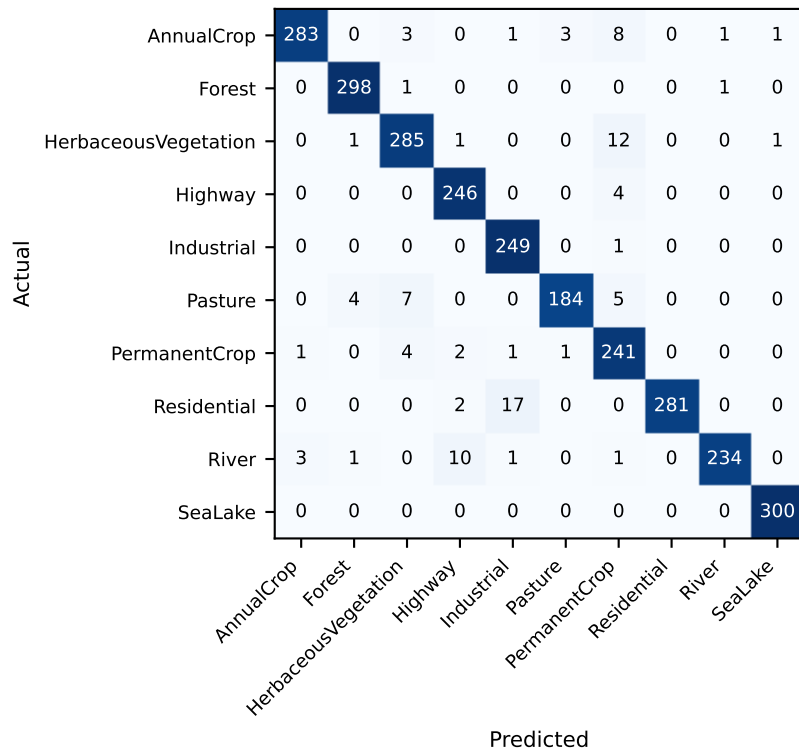
AUC (macro, OVR)



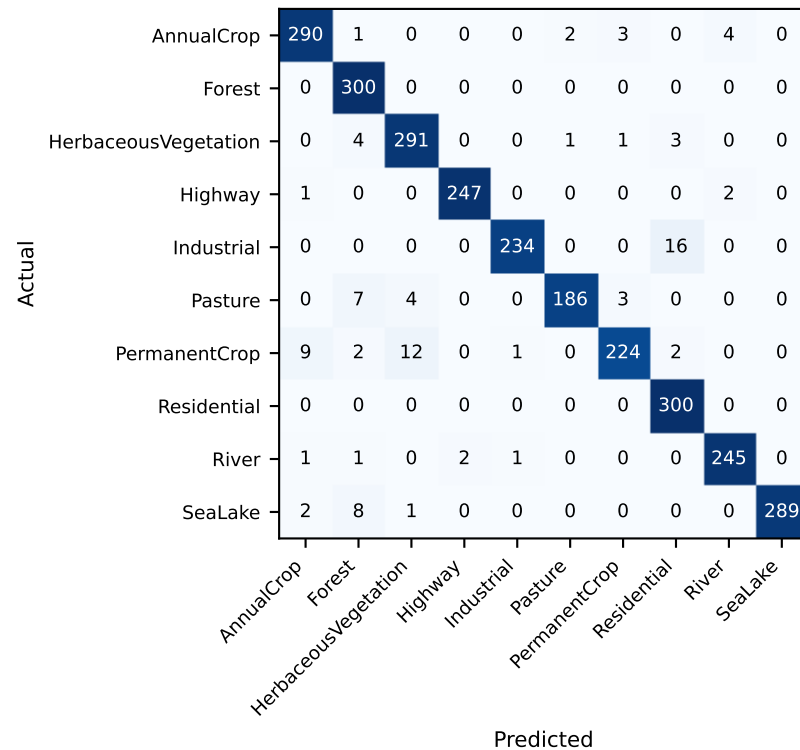
Avg. Loss



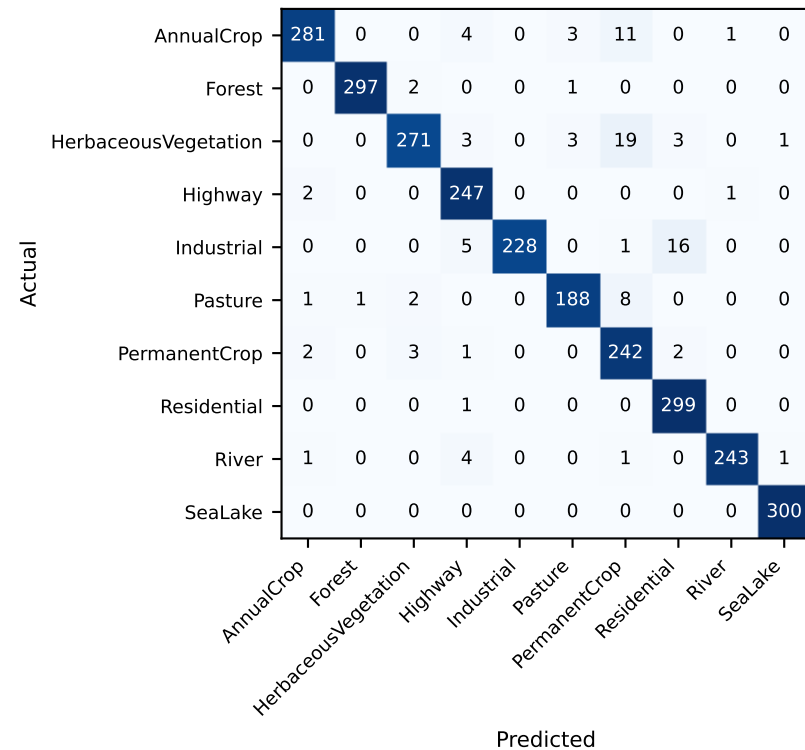
ConvNeXt



EfficientNetV2



ResNeXt



AUTOMATED CONCLUSIONS

=====

Overall pick: EfficientNetV2 – best average rank across 7 compared metrics, and clearly ahead of ConvNeXt (by 3%+) on 1 of them vs. 0 the other way. Average rank (0 = best): EfficientNetV2 0.00, ConvNeXt 1.14, ResNeXt 1.86.

[ConvNeXt]

Strengths: none stand out

Weaknesses: none stand out

[EfficientNetV2]

Strengths: Avg. Loss (0.0997)

Weaknesses: none stand out

[ResNeXt]

Strengths: none stand out

Weaknesses: none stand out

(Strength = best on a metric by 3%+ over the runner-up; weakness = worst by 3%+ under the second-worst. Shown but not counted: Balanced Accuracy – it duplicates macro recall.)